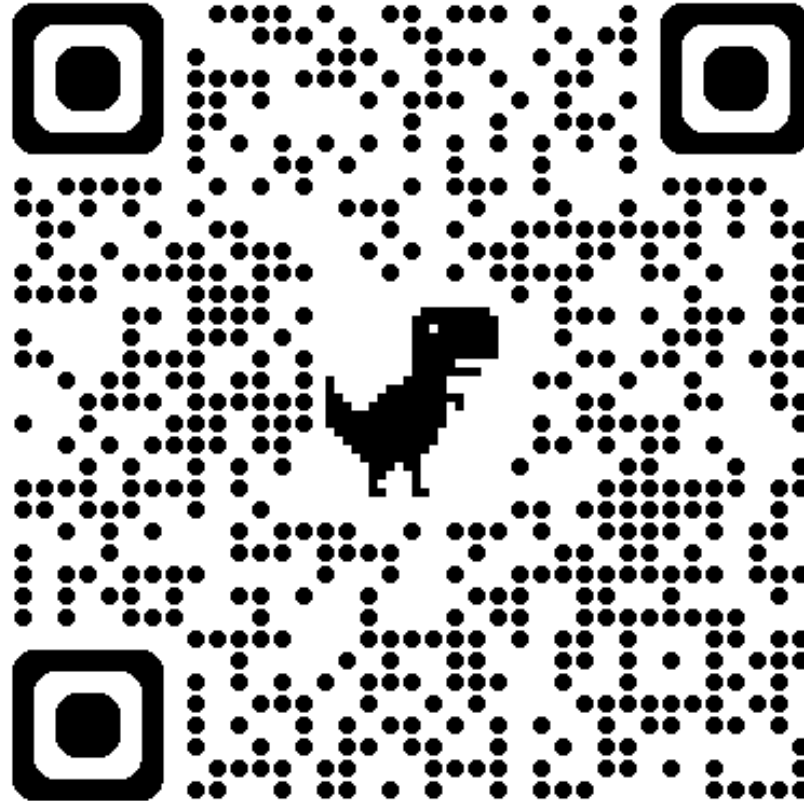




Introduction to GPU Acceleration

View the Slides



https://github.com/ResearchComputing/Intro_GPU_Acceleration

Meet the User Support Team



Layla
Freeborn



Brandon
Reyes



Andy
Monaghan



Michael
Schneider



John
Reiland



Dylan
Gottlieb





Mohal
Khandelwal

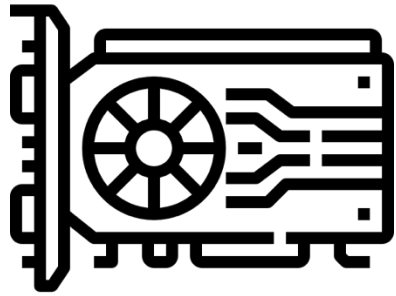


Ragan
Lee

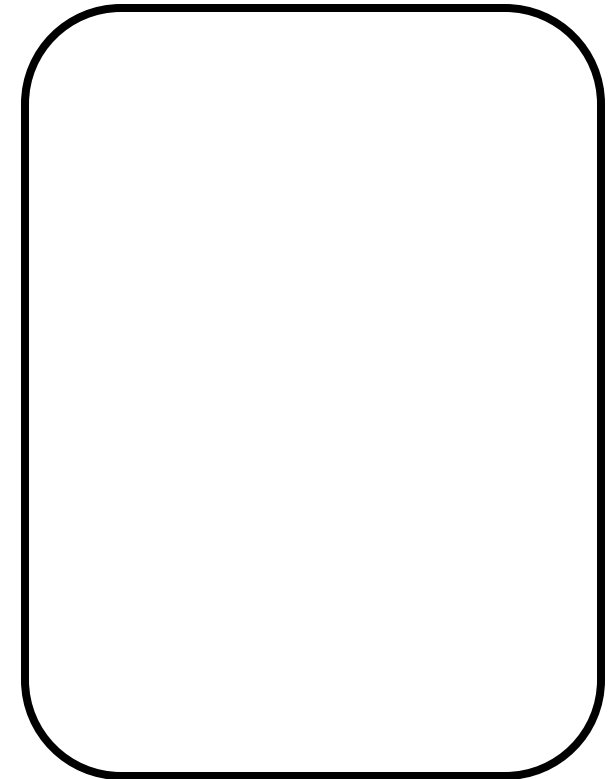
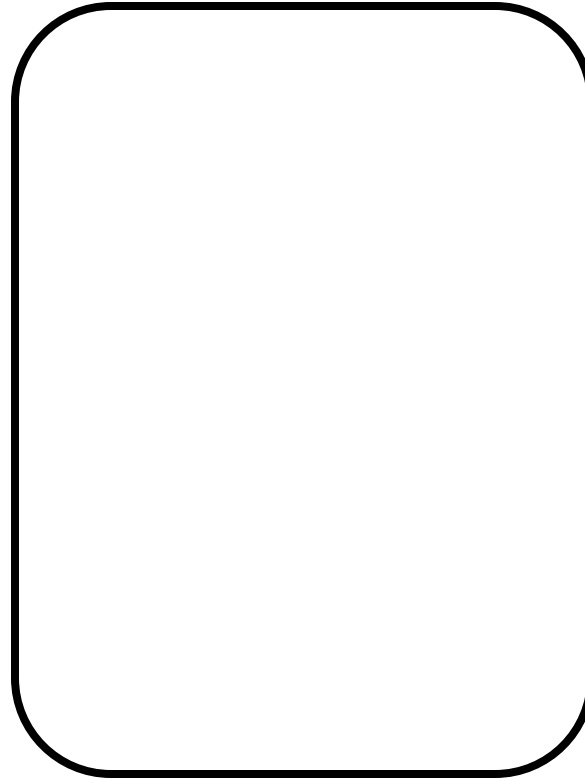
Session Overview



Session Overview

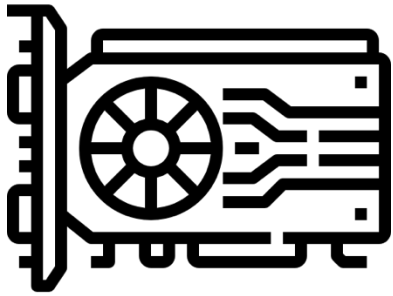


**Basics
Of GPUs**

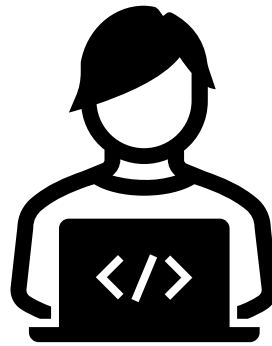


[GPU Icon](#)

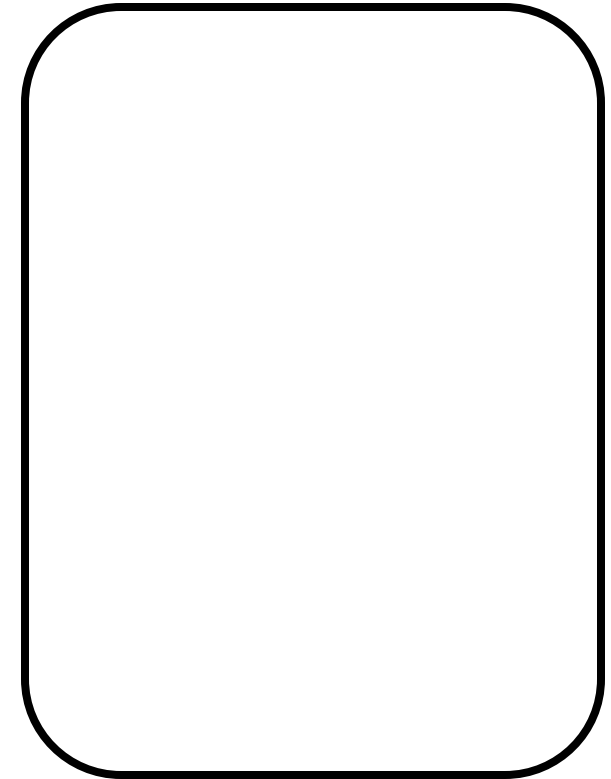
Session Overview



**Basics
Of GPUs**

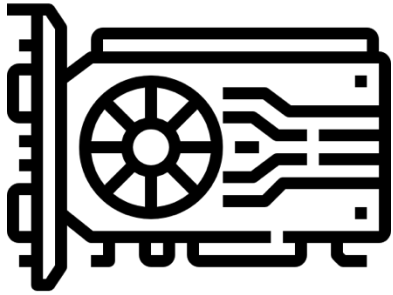


**Code
Optimization**

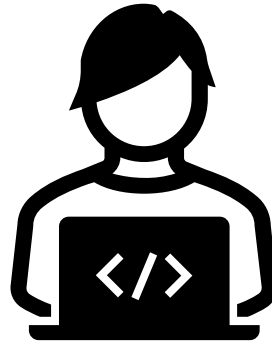


[GPU Icon](#)


Session Overview



**Basics
Of GPUs**

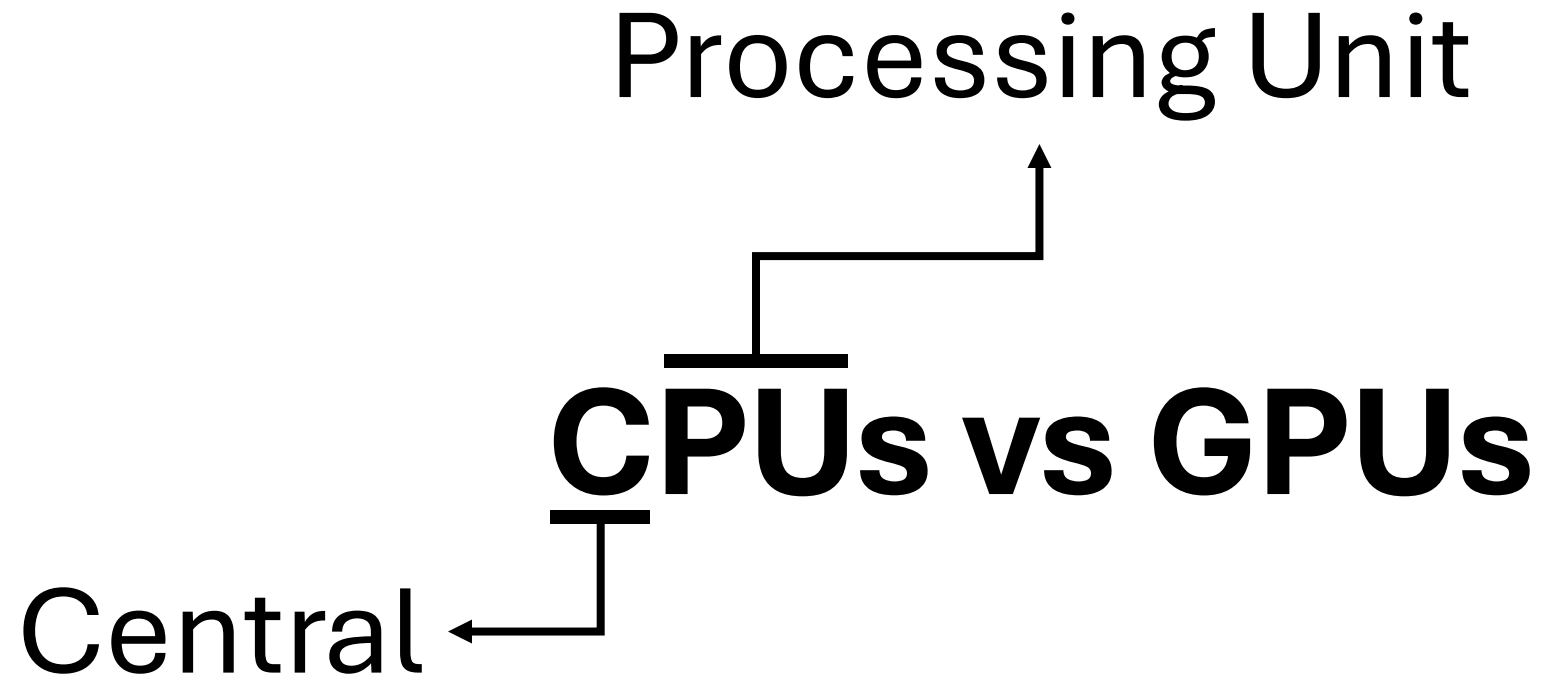


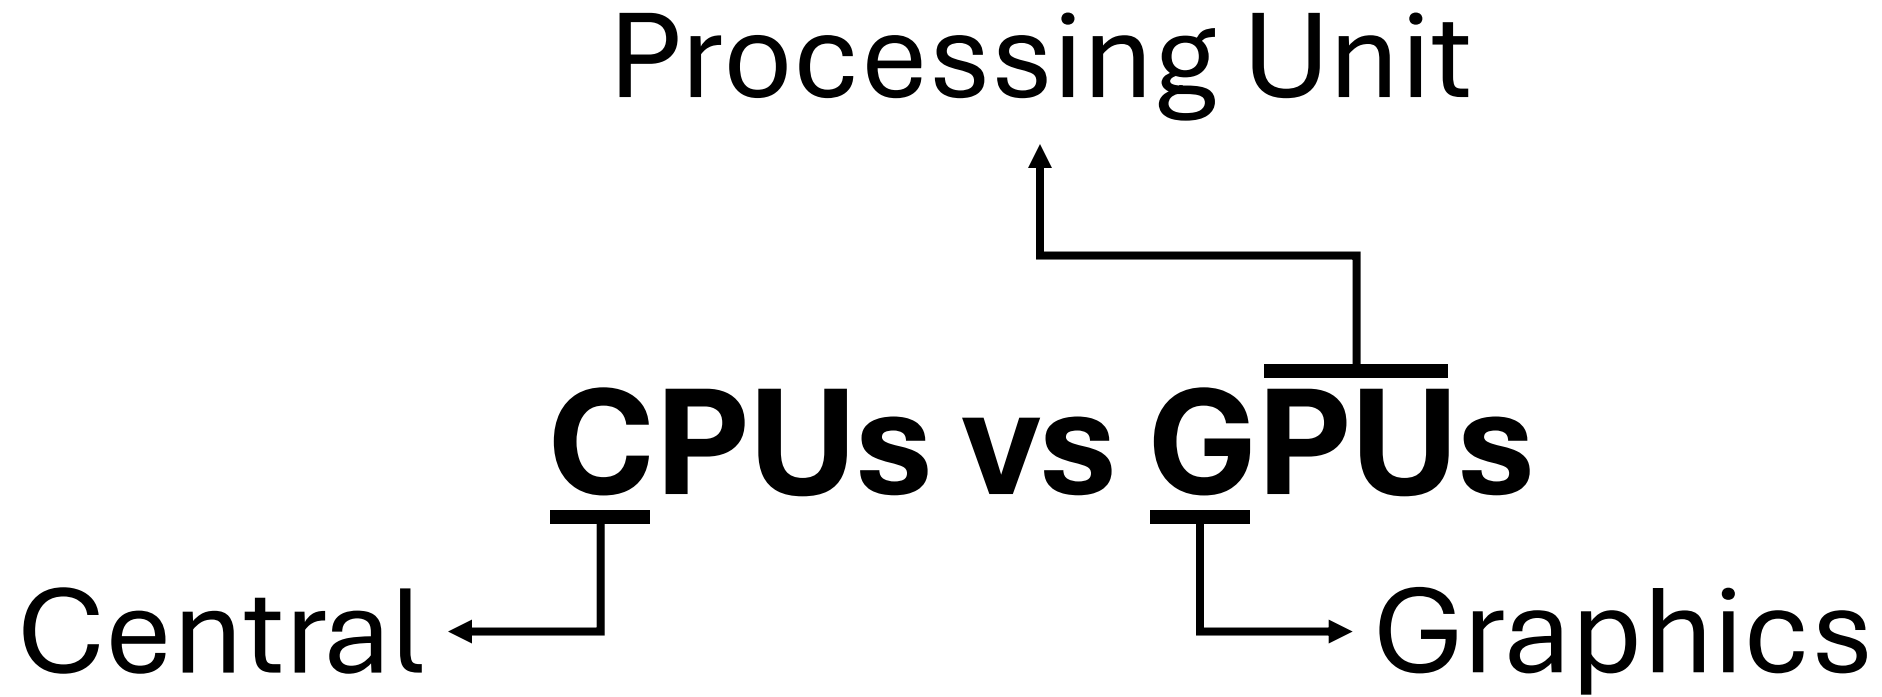
**Code
Optimization**




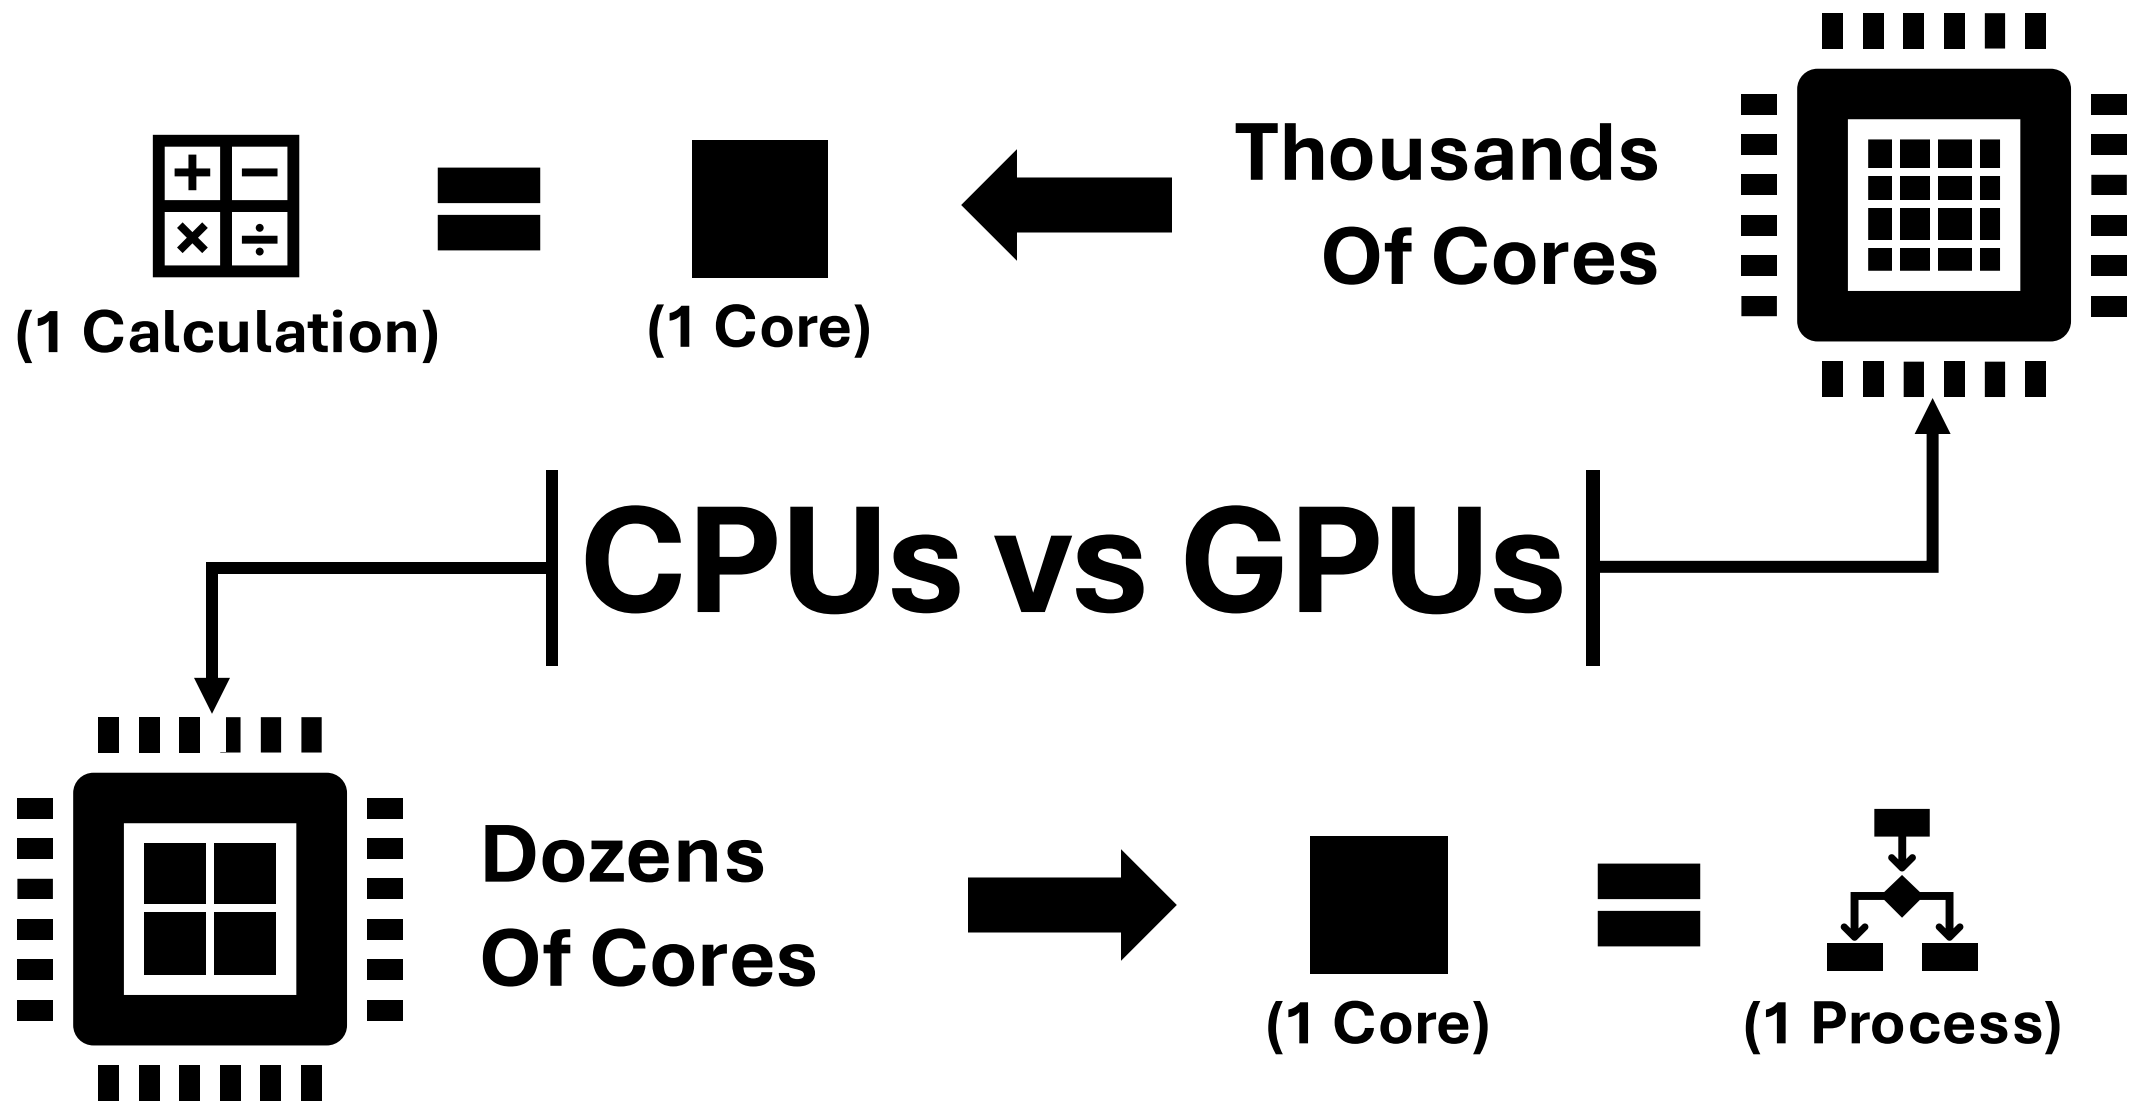
**Monitoring
GPU Usage**

CPU vs GPU

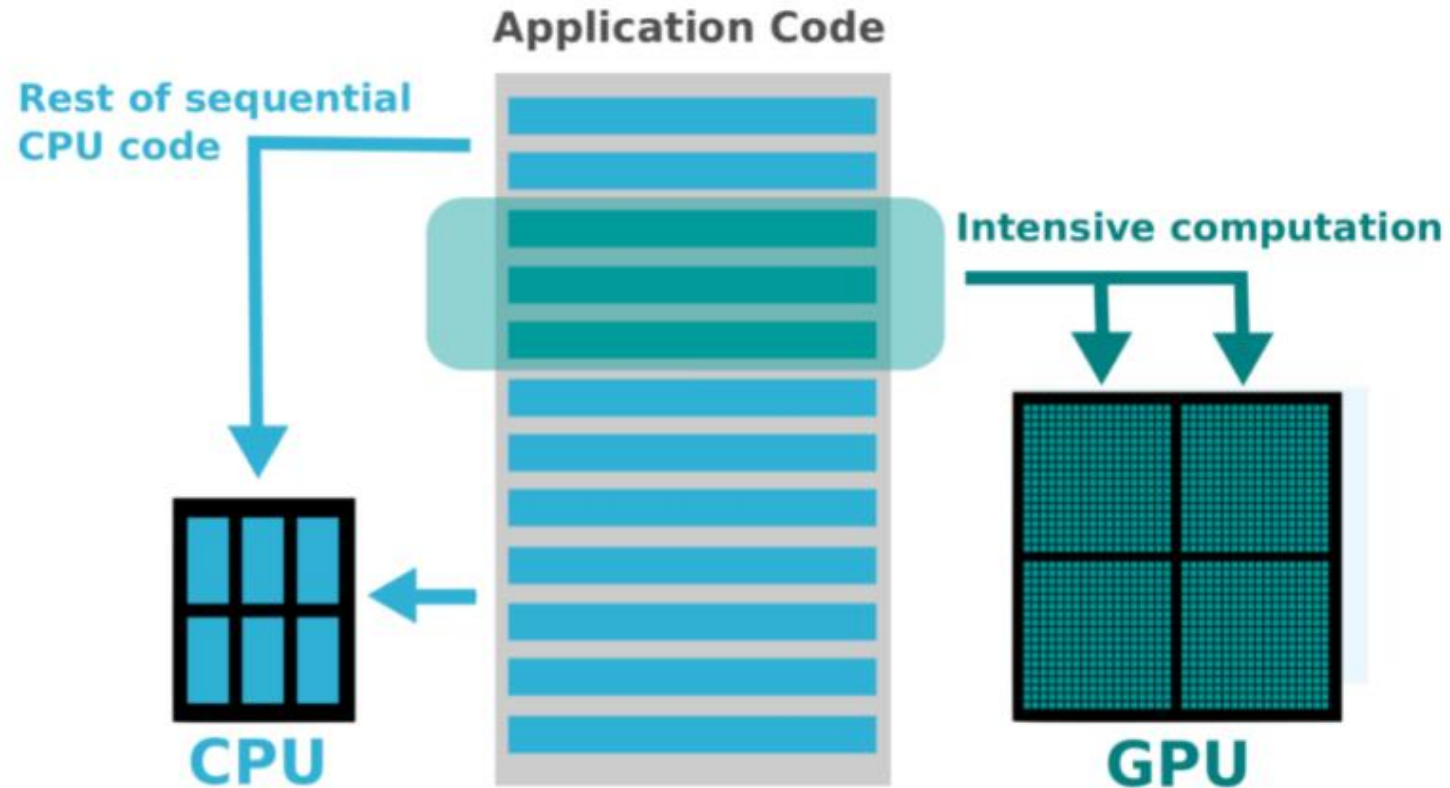






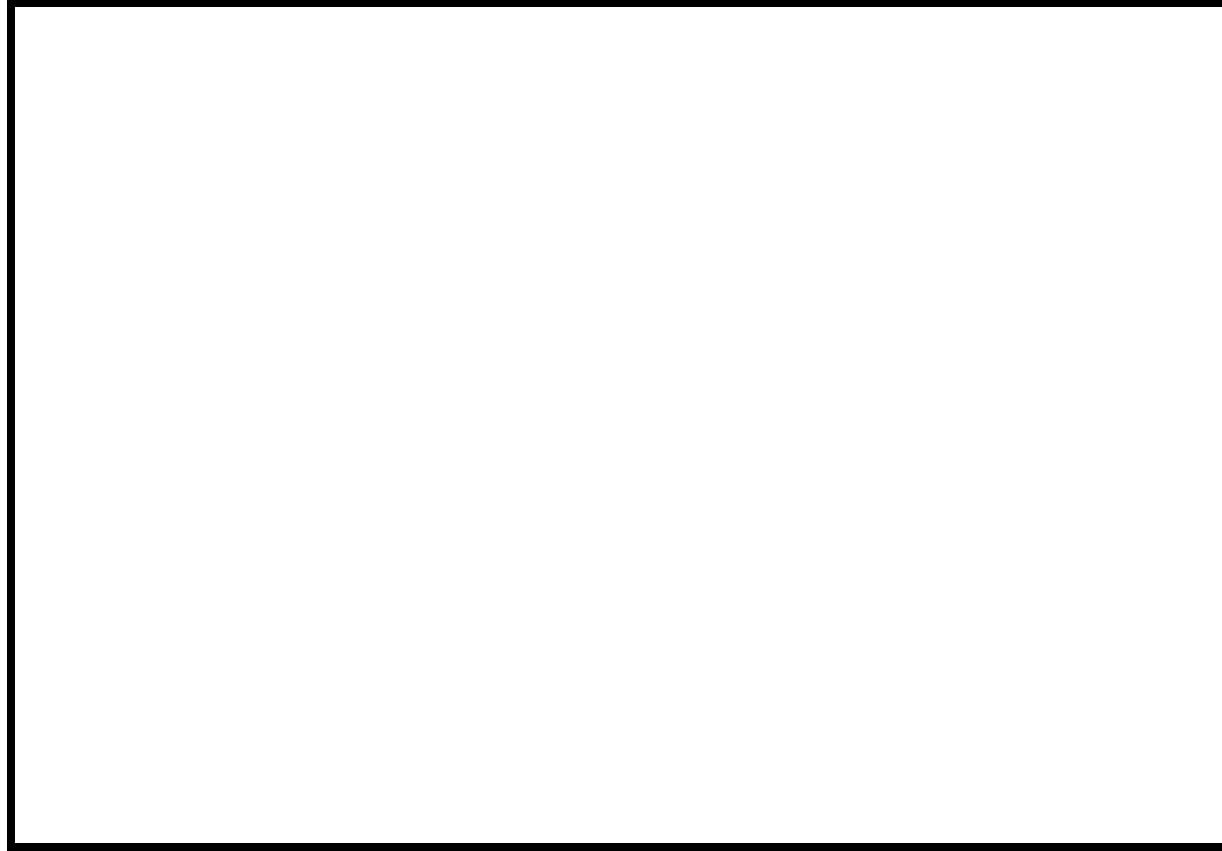


Computational Offloading



[Graphic Source](#)

GPU Acceleration Checklist

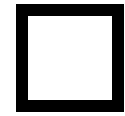


GPU Acceleration Checklist

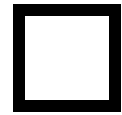


Computational Intensity

GPU Acceleration Checklist



Computational Intensity



Algorithmic Complexity

GPU Acceleration Checklist

☐ **Computational Intensity**

☐ **Algorithmic Complexity**

☐ **Data Type**

GPU Acceleration Checklist

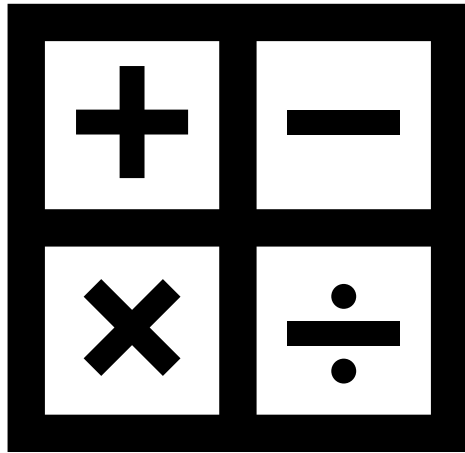
☐ **Computational Intensity**

☐ **Algorithmic Complexity**

☐ **Data Type**

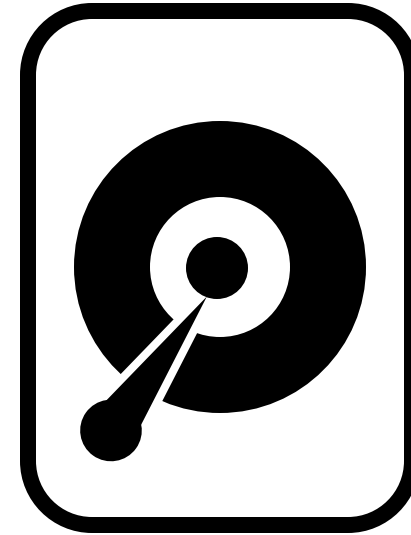
☐ **Data Dependency**

Computational Intensity



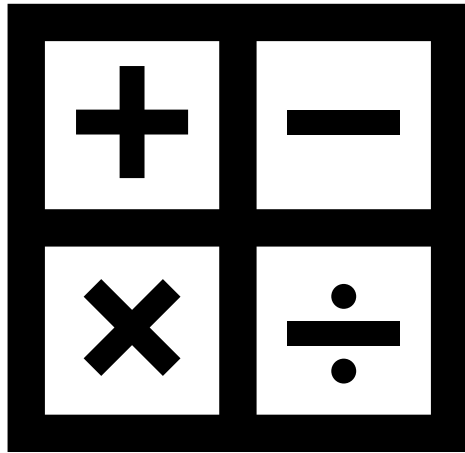
Calculations

VS



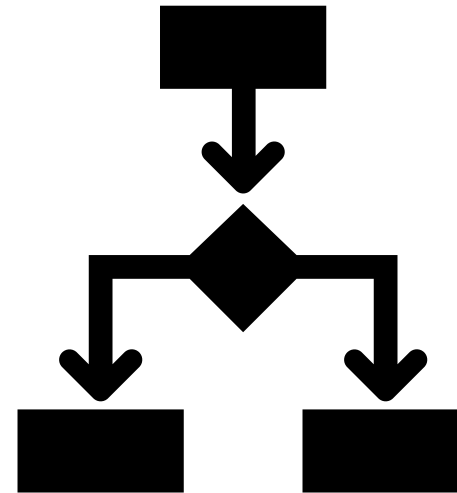
Data Access

Algorithmic Complexity



Calculations

VS



Branching Logic

Data Type

123

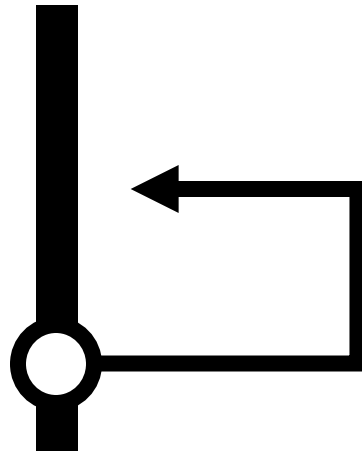
Numeric

vs

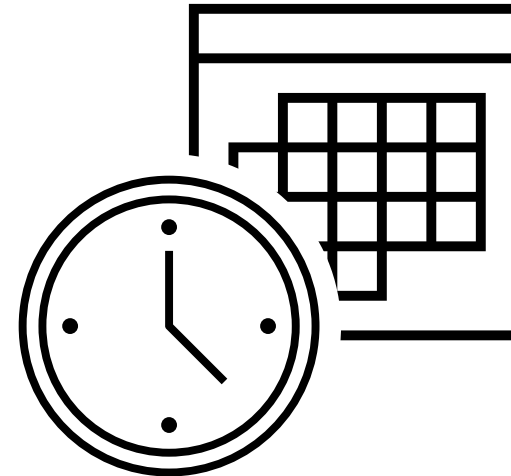
“Text”

Complex Objects

Data Dependency



Recursion



**Temporal
Time / Dates**

GPU Acceleration Checklist

☐ **Computational Intensity**

☐ **Algorithmic Complexity**

☐ **Data Type**

☐ **Data Dependency**

Alpine GPUs

	NVIDIA			AMD
Type	A100	L40	GH200	MI100
Cores	7k	15K	17k	7.7k
VRAM	40 / 80	48	96	32
Purpose	General	Viz, AI Inference	AI Training, High Data I/O	Scientific

Alpine GPUs

	NVIDIA			AMD
Type	A100	L40	GH200	MI100
Partition	aa100	al40	gh200*	ami100
Nodes	40 (8) / 80 (4)	3	2	2
GPUs Per Node	3	3	1	3

Requesting GPUs

SLURM Directives:

--partition= < >

--gres=gpu:<#>

--ntasks=<#>

Requesting GPUs

SLURM Directives:

--partition= < >

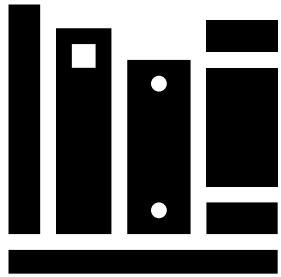
--gres=gpu:<#>

--ntasks=<#>

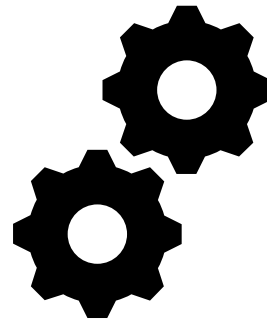
sinteractive --partition=ami100 --gres=gpu:2 --ntasks=20

#SBATCH <directive>


Code Optimization



Libraries



Directives



Languages

Key Terms

- Host == CPU
- Device == GPU
- Kernel == Functions launched on GPU

GPU Libraries – Drop in Replacement

#create dataset with 100,000 points

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```


#run DBSCAN clustering algorithm

```
from sklearn.cluster import DBSCAN
```


```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```

Dataset



Clusters



GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
from sklearn.datasets import make_circles
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#run DBSCAN clustering algorithm
from sklearn.cluster import DBSCAN
db = DBSCAN(eps=0.6, min_samples=2)
y_db = db.fit_predict(X)
```

GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points  
from sklearn.datasets import make_circles  
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#convert dataset to Pandas DataFrame
```

```
#run DBSCAN clustering algorithm  
from sklearn.cluster import DBSCAN  
db = DBSCAN(eps=0.6, min_samples=2)  
y_db = db.fit_predict(X)
```

GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
from sklearn.datasets import make_circles
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)

#convert dataset to Pandas DataFrame
import pandas as pd
import cudf
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
X_gpu = cudf.DataFrame.from_pandas(X_df)

#run DBSCAN clustering algorithm
from sklearn.cluster import DBSCAN
db = DBSCAN(eps=0.6, min_samples=2)
y_db = db.fit_predict(X)
```

GPU Libraries – Drop in Replacement

#create dataset with 100,000 points

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

#convert dataset to Pandas DataFrame

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

#run DBSCAN clustering algorithm

```
from sklearn.cluster import DBSCAN
```

```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```



#run GPU-accelerated DBSCAN

```
from cuml import DBSCAN
```

GPU Libraries – Drop in Replacement

#create dataset with 100,000 points

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

#convert dataset to Pandas DataFrame

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

~~#run DBSCAN clustering algorithm~~

~~from sklearn.cluster import DBSCAN~~

~~db = DBSCAN(eps=0.6, min_samples=2)~~

~~y_db = db.fit_predict(X)~~



#run GPU-accelerated DBSCAN

```
from cuml import DBSCAN
```

GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
```

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#convert dataset to Pandas DataFrame
```

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

```
#run GPU-accelerated DBSCAN
```

```
from cuml import DBSCAN
```

```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```


GPU-Enabled Frameworks



GPU Compiler Directives

GPU Compiler Directives

Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

GPU Compiler Directives

Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

GPU Compiler Directives

Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

Data directives

Generate code to manage specific data operations to support parallelism

GPU Compiler Directives

Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

Data directives

Generate code to manage specific data operations to support parallelism

```
//Hello_World_OpenACC.c
#pragma acc data copy(a)
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

Languages

- OpenCL (NVIDIA, AMD, & CPUs)
 - Flexible / portable option
- HIP (AMD -> NVIDIA)
 - AMD developed
 - Can convert CUDA code via `hippify`
- CUDA (NVIDIA only)
 - Most robust and largest developer community

Monitoring GPU Usage

- Nvidia-smi
- rocm-smi

NVIDIA-SMI 510.47.03				Driver Version: 510.47.03		CUDA Version: 11.6	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M. MIG M.
0	NVIDIA A100-PCI...	Off	00000000:21:00.0	Off	0	Default	Disabled
N/A	36C	P0	40W / 250W	0MiB / 40960MiB	0%	Default	Disabled
1	NVIDIA A100-PCI...	Off	00000000:81:00.0	Off	0	Default	Disabled
N/A	36C	P0	40W / 250W	0MiB / 40960MiB	0%	Default	Disabled
2	NVIDIA A100-PCI...	Off	00000000:E2:00.0	Off	0	Default	Disabled
N/A	37C	P0	40W / 250W	0MiB / 40960MiB	0%	Default	Disabled
Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	Usage
	ID	ID					
No running processes found							

Troubleshooting GPU Workflows

- Is your application and/or code GPU accelerated?

Confirm that you installed the GPU accelerated version!

- Does your application or code support **multi**-GPU acceleration?
- Is your application ROCM- or CUDA-aware?

You can't run CUDA code on AMD GPUs. Not all applications are available for AMD GPUs.

- Can your application “see” the GPU?
- Did you request enough CPUs and RAM?

Documentation



<https://curc.readthedocs.io/en/latest/>

Survey and feedback



Survey: <http://tinyurl.com/curc-survey18>