




# Introduction to GPU Acceleration

March 6, 2026

Andrew Monaghan

[rc-help@colorado.edu](mailto:rc-help@colorado.edu)

# View the Slides



[https://github.com/ResearchComputing/Intro\\_GPU\\_Acceleration](https://github.com/ResearchComputing/Intro_GPU_Acceleration)



# Meet the User Support Team



Layla  
Freeborn



Brandon  
Reyes



Andy  
Monaghan



Michael  
Schneider



John  
Reiland



Dylan  
Gottlieb




Mohal  
Khandelwal




Ragan  
Lee

# Session Overview




**Basics  
Of GPUs**

**1**



**Code  
Optimization**

**2**





**Monitoring  
GPU Usage**


**3**

[GPU Icon](#)


# CPU vs GPU













# Computational Offloading



[Graphic Source](#)

# Data Offloading CPU -> GPU



Data needs to be copied from CPU to GPU, computation is on the GPU, then output is transferred back to CPU.

# Criteria for GPU Acceleration

1. The time spent on computationally intensive parts of the workflow exceeds the time spent transferring data to and from GPU memory
2. Computations are massively parallel- the computations can be broken down into hundreds or thousands of independent units of work



# Factors Affecting GPU Speedup

- 1 Computational Intensity
- 2 Data Dependency
- 3 Data Type
- 4 Code/Algorithmic Complexity

# Factors Affecting GPU Speedup

1

## Computational Intensity

GPUs perform best when there is a lot of processing compared to loading and storing data (FLOP per Byte ratio).

# Factors Affecting GPU Speedup

## 2 Data Dependency

**Data Dependency-** A situation in which an instruction is dependent on a result from a sequentially previous instruction before it can complete its execution. (avoid!)

# Factors Affecting GPU Speedup

## 3 Data Type

- Operations on strings are slow unless they can be treated as numbers.
- Performance per GPU can vary if workflows include 16-bit and 64-bit floats.

# Factors Affecting GPU Speedup

4

## Code/Algorithmic Complexity

Simple code is better ported to GPUs.

- Deeply-branched code and while-loops may perform poorly on GPUs
- Recursive functions need to be re-written

# Alpine GPU Partitions

|           | NVIDIA                   |              |                               | AMD                        |
|-----------|--------------------------|--------------|-------------------------------|----------------------------|
| Partition | aa100<br>(atesting_a100) | al40         | gh200*                        | ami100<br>(atesting_mi100) |
| # Nodes   | 12                       | 3            | 2                             | 8                          |
| GPU Type  | A100                     | L40          | GH200                         | MI100                      |
| GPUs/Node | 3                        | 3            | 1                             | 3                          |
| Cores/GPU | 7k                       | 15K          | 17k                           | 7.7k                       |
| VRAM/GPU  | 40 / 80                  | 48           | 96                            | 32                         |
| Purpose   | General                  | AI Inference | AI Training,<br>High Data I/O | General                    |



# Exploring Alpine GPU Partitions

Try these commands.

```
$ ssh <username>@login.rc.colorado.edu
```

```
$ sinfo --Format Partition
```

```
$ sinfo --partition aa100,ami100,atesting_a100,atesting_mi100 --Format=Partition,Nodes,Time
```

```
$ scontrol show partition atesting_a100
```

```
$ scontrol show node c3gpu-c2-u13
```

# Requesting Alpine GPUs with Slurm batch script

Slurm flags needed to request 1 NVIDIA GPU node with 2 GPUs and 20 CPU cores

```
--partition=aa100  
--gres=gpu:2  
--ntasks=20
```

in a job script  
submitted with  
***sbatch***  
command

```
#SBATCH --partition=aa100  
#SBATCH --qos=normal  
#SBATCH --gres=gpu:2  
#SBATCH --nodes=1  
#SBATCH --ntasks=20  
#SBATCH --time=12:00:00  
#SBATCH --job-name=gpu_test  
#SBATCH --output=gpu_test_%j.out  
#SBATCH --error=gpu_test_%j.err
```


# Requesting Alpine GPUs with Slurm interactive

```
# request one NVIDIA GPU on the A100 testing partition:
```

```
$ sinteractive --partition=atesting_a100 --qos=testing --gres=gpu:1 --nodes=1 --ntasks=10 --  
time=1:00:00
```


*Note: Queue waits for non-testing GPU partitions are typically long (12-24 hours). Therefore, interactive jobs -- which require you to wait for the job to start -- will usually be run on the testing partitions during the onboarding and debugging stages of your workflow.*

# Code Optimization



[GPU Icon](#)

# GPU Libraries – Drop in Replacement



The screenshot shows the RAPIDS website homepage. The header is purple with white text for navigation: RAPIDS, Ecosystem, Learn More, News, User Guides, Install, Docs, and GitHub. The main hero section is purple with the text 'RAPIDS GPU Accelerated Data Science' and a 'QUICK START' button. Below this is a row of five news items, each with a bell icon and a title: 'NEW: CUML NOW ACCELERATES SCIKIT-LEARN', 'NEW: POLARS GPU ENGINE', 'CUDF NOW PRE-INSTALLED IN GOOGLE COLAB', 'VECTOR SEARCH NOW WITH CUVS', and 'RAPIDS 25.10 RELEASED'. The footer contains three columns: 'WHAT IS RAPIDS' (describing unmatched speed with familiar APIs), 'WHY USE RAPIDS' (describing fluid interaction with data), and 'OPEN SOURCE ECOSYSTEM' (stating it's open source and available on GitHub).

RAPIDS Ecosystem Learn More News User Guides Install Docs GitHub

## RAPIDS

GPU Accelerated Data Science

QUICK START

NEW: CUML NOW ACCELERATES SCIKIT-LEARN NEW: POLARS GPU ENGINE CUDF NOW PRE-INSTALLED IN GOOGLE COLAB VECTOR SEARCH NOW WITH CUVS RAPIDS 25.10 RELEASED

### WHAT IS RAPIDS

RAPIDS provides unmatched speed with familiar APIs that match the most popular PyData libraries. Built on state-of-the-art foundations like [NVIDIA CUDA](#) and [Apache Arrow](#), it unlocks the speed of GPUs with code you already know. [Jump to About Section](#)

### WHY USE RAPIDS

RAPIDS allows fluid, creative interaction with data for everyone from BI users to AI researchers on the cutting edge. GPU acceleration means less time and less cost moving data and training models. [Jump to RAPIDS Use Cases](#)

### OPEN SOURCE ECOSYSTEM


RAPIDS is Open Source and available on [GitHub](#). Our mission is to empower and advance the open-source GPU data science data engineering ecosystem. [Jump to RAPIDS GitHub](#)

<https://rapids.ai>

# GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
from sklearn.datasets import make_circles
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#run DBSCAN clustering algorithm
from sklearn.cluster import DBSCAN
db = DBSCAN(eps=0.6, min_samples=2)
y_db = db.fit_predict(X)
```



# GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
from sklearn.datasets import make_circles
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#run DBSCAN clustering algorithm
from sklearn.cluster import DBSCAN
db = DBSCAN(eps=0.6, min_samples=2)
y_db = db.fit_predict(X)
```



# GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points  
from sklearn.datasets import make_circles  
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#convert dataset to Pandas DataFrame
```

```
#run DBSCAN clustering algorithm  
from sklearn.cluster import DBSCAN  
db = DBSCAN(eps=0.6, min_samples=2)  
y_db = db.fit_predict(X)
```

# GPU Libraries – Drop in Replacement

```
#create dataset with 100,000 points
```

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

```
#convert dataset to Pandas DataFrame
```

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

```
#run DBSCAN clustering algorithm
```

```
from sklearn.cluster import DBSCAN
```

```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```

# GPU Libraries – Drop in Replacement

**#create dataset with 100,000 points**

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

**#convert dataset to Pandas DataFrame**

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

**#run DBSCAN clustering algorithm**

```
from sklearn.cluster import DBSCAN
```

```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```



**#run GPU-accelerated DBSCAN**

```
from cuml import DBSCAN
```

# GPU Libraries – Drop in Replacement

**#create dataset with 100,000 points**

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

**#convert dataset to Pandas DataFrame**

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

~~#run DBSCAN clustering algorithm~~

~~from sklearn.cluster import DBSCAN~~

~~db = DBSCAN(eps=0.6, min\_samples=2)~~

~~y\_db = db.fit\_predict(X)~~



**#run GPU-accelerated DBSCAN**

```
from cuml import DBSCAN
```

# GPU Libraries – Drop in Replacement

**#create dataset with 100,000 points**

```
from sklearn.datasets import make_circles
```

```
X, y = make_circles(n_samples=int(1e5), factor=.35, noise=.05)
```

**#convert dataset to Pandas DataFrame**

```
import pandas as pd
```

```
import cudf
```

```
X_df = pd.DataFrame({'fea%d'%i: X[:,i] for i in range(X.shape[1])})
```

```
X_gpu = cudf.DataFrame.from_pandas(X_df)
```

**#run GPU-accelerated DBSCAN**

```
from cuml import DBSCAN
```

```
db = DBSCAN(eps=0.6, min_samples=2)
```

```
y_db = db.fit_predict(X)
```

# Hands on exercise! Compare DBSCAN cpu vs gpu

Go back to the interactive job you started earlier in the terminal. Do the following:

```
# go to your scratch directory
```

```
$ cd /scratch/alpine/$USER
```

```
# clone the git repository for this course to get the test code
```

```
$ git clone https://github.com/ResearchComputing/Intro_GPU_Acceleration
```

```
# go into your git repository
```

```
$ cd Intro_GPU_Acceleration
```

# Hands on continued...

```
# now load a python environment that has NVIDIA RAPIDS  
$ module load miniforge  
$ mamba activate /curc/sw/conda_env/rapids-25.10
```

```
# run the CPU version of the DBSCAN code  
$ python dbscan_cpu.py
```

```
# run the GPU version of the DBSCAN code  
$ python dbscan_gpu.py
```

...what magnitude difference is there between the CPU and GPU trials?

# GPU-Enabled Frameworks (deep learning)



Known for: comprehensiveness, scalability



Known for: Flexibility, pythonic, intuitive



Known for: Ease of use



Known for: Distributed training



## GPU Compiler Directives

### Key Terms

- Host == CPU
- Device == GPU
- Kernel == Functions launched on GPU

## GPU Compiler Directives

### **Kernel directives**

Generate parallel accelerator kernels  
for the loop following the directive.

# GPU Compiler Directives

## Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

# GPU Compiler Directives

## Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

## Data directives

Generate code to manage specific data operations to support parallelism

# GPU Compiler Directives

## Kernel directives

Generate parallel accelerator kernels for the loop following the directive.

```
//Hello_World_OpenACC.c
void Print_Hello_World()
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

## Data directives

Generate code to manage specific data operations to support parallelism

```
//Hello_World_OpenACC.c
#pragma acc data copy(a)
{
    #pragma acc kernels
    for(int i=0; i<5; i++)
    {
        printf("Hello World!\n")
    }
}
```

# GPU Languages

- OpenCL (NVIDIA, AMD, & CPUs)
  - Flexible / portable option
- HIP (AMD <-> NVIDIA)
  - AMD developed
  - Can convert CUDA code via `hippify`
- CUDA (NVIDIA only)
  - Most robust and largest developer community

# Monitoring GPU Usage

```
$ nvidia-smi
```


```
$ rocm-smi
```

```
$ seff
```

[More info & limitations.](#)

|                                                                             |        |               |               |                  |                 |            |         |          |  |
|-----------------------------------------------------------------------------|--------|---------------|---------------|------------------|-----------------|------------|---------|----------|--|
| NVIDIA-SMI 510.47.03      Driver Version: 510.47.03      CUDA Version: 11.6 |        |               |               |                  |                 |            |         |          |  |
| GPU                                                                         | Name   | Persistence-M |               | Bus-Id           | Disp.A          | Volatile   | Uncorr. | ECC      |  |
| Fan                                                                         | Temp   | Perf          | Pwr:Usage/Cap |                  | Memory-Usage    | GPU-Util   | Compute | M.       |  |
|                                                                             |        |               |               |                  |                 |            | MIG     | M.       |  |
| 0                                                                           | NVIDIA | A100-PCI...   | Off           | 00000000:21:00.0 | Off             |            |         | 0        |  |
| N/A                                                                         | 36C    | P0            | 40W / 250W    |                  | 0MiB / 40960MiB | 0%         | Default | Disabled |  |
| 1                                                                           | NVIDIA | A100-PCI...   | Off           | 00000000:81:00.0 | Off             |            |         | 0        |  |
| N/A                                                                         | 36C    | P0            | 40W / 250W    |                  | 0MiB / 40960MiB | 0%         | Default | Disabled |  |
| 2                                                                           | NVIDIA | A100-PCI...   | Off           | 00000000:E2:00.0 | Off             |            |         | 0        |  |
| N/A                                                                         | 37C    | P0            | 40W / 250W    |                  | 0MiB / 40960MiB | 0%         | Default | Disabled |  |
| Processes:                                                                  |        |               |               |                  |                 |            |         |          |  |
| GPU                                                                         | GI     | CI            | PID           | Type             | Process name    | GPU Memory | Usage   |          |  |
|                                                                             | ID     | ID            |               |                  |                 |            |         |          |  |
| No running processes found                                                  |        |               |               |                  |                 |            |         |          |  |

# Monitoring GPU Usage on NVIDIA GPUs



The screenshot shows a web page from the University of Colorado Boulder Research Computing department. The page is titled "Profiling NVIDIA GPU Performance". It includes a sidebar with navigation links such as "Getting Started", "The Compute Environment", "Clusters", "Portals & Gateways", and "Running Jobs". The main content area explains that NVIDIA Performance Counters provide low-level metrics on GPU usage and lists three tools: `nvidia-smi`, `Nsight Compute (ncu)`, and `Nsight Systems (nsys)`. Below this, it provides "Sample CUDA Code" with two examples: "Vector Addition" and "Matrix Multiplication". The "Matrix Multiplication" example is selected, showing a CUDA program that adds two vectors of floats.

**Profiling NVIDIA GPU Performance**

The NVIDIA Performance Counters provide low-level metrics on GPU usage, enabling users to understand how efficiently their code uses the GPU. This is especially important for optimizing workloads on Alpine's A100 GPU nodes, where GPU time is a valuable and shared resource.

The following tools are available for interacting with performance counters:

- `nvidia-smi`: For basic monitoring of GPU resource usage.
- `Nsight Compute (ncu)`: For detailed GPU kernel performance analysis.
- `Nsight Systems (nsys)`: For system-wide GPU and CPU performance tracing.

**Sample CUDA Code**

To demonstrate how to use NVIDIA profiling and monitoring tools, we will use two different CUDA programs:

- Vector Addition (`vectorAdd.cu`): A simpler operation that serves as a quick example for profiling.
- Matrix Multiplication (`matrixMultiply.cu`): A compute-intensive task designed to showcase GPU load during complex operations.

`vectorAdd.cu` `matrixMultiply.cu`

Here's a CUDA program `vectorAdd.cu`, that adds two vectors of floats.

```
#include <iostream>
#include <cuda_runtime.h>

#define N 1024 // Size of the vectors

__global__ void vectorAdd(float *A, float *B, float *C) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < N) {
        C[idx] = A[idx] + B[idx];
    }
}

int main() {
```

<https://curc.readthedocs.io/en/latest/programming/profiling-nvidia-gpu-performance.html>



# Troubleshooting GPU Workflows

- Is your application and/or code GPU accelerated?

*Confirm that you installed the GPU accelerated version!*

- Does your application or code support **multi**-GPU acceleration?
- Is your application ROCM- or CUDA-aware?

*You can't run CUDA code on AMD GPUs. Not all applications are available for AMD GPUs.*

- Can your application “see” the GPU?
- Did you request enough CPUs and RAM?

# Documentation



<https://curc.readthedocs.io/en/latest/>

# Survey and feedback



Survey: <http://tinyurl.com/curc-survey18>